

---

# Symbolic Mutation Compression and RAM-Optimized Buffer Layout in EchoPulse

## 1. Introduction

The EchoPulse Key Encapsulation Mechanism (KEM) relies on a deterministic graph mutation function  $\mu(G,t)$  to evolve the underlying symbolic graph  $G_t$ . While the previous "Low-RAM Operating Mode" document addressed overall graph storage in Flash/ROM, this document focuses specifically on optimizing the RAM footprint of the *mutation result* itself and its associated processing buffer. The goal is to define a low-footprint mutation representation that can be efficiently generated, applied, and fully reconstructed deterministically within extremely constrained RAM environments (e.g.,  $\leq 512$  bytes for mutation-specific data), preserving the integrity and unpredictability properties critical to EchoPulse's security.

## 2. Symbolic Compression Schemes for Mutation Results (Symbolic Compression Engineer)

The mutation function  $\mu(G,t)$  transforms the base graph  $G_{base}$  into  $G_t$  by altering a subset of its  $\delta(v,s) \rightarrow v'$  mappings. Instead of storing a full  $\delta_t$  table in RAM (which is prohibitive), we store a compressed representation of the *changes* or *rules* applied by  $\mu$ .

- **Core Idea:**  $\mu$  produces a "patch" or a "set of instructions" that, when applied to  $G_{base}$ , yields  $G_t$ . This patch is what needs to be compressed.
- **Proposed Bit-Level Encoding Schemes:**
  1. **Symbol Remapping via Delta Tables / XOR Diffs:**
    - **Concept:** For a given state  $v$ , the mutation might permute the next states for all symbols  $s \in \Sigma$ . Instead of storing explicit new target states, store a permutation vector or a bitmask/XOR difference applied to the base state.
    - **Example:** If  $\delta_{base}(v,s) = s \oplus C_v$  for some constant  $C_v$ ,  $\mu$  could modify  $C_v$  or apply a permutation  $\pi_t$  to  $s$ . The mutation result could be  $\pi_t$  itself or parameters defining it.
    - **Format:** For each affected state  $v$ , store a bitmask or a small index into a pre-defined set of permutations/modifications.
    - **Size Implications:** If  $\mu$  affects  $k$  states, and for each state it applies a chosen permutation from a library of  $P$  permutations, this could be  $k \times \log_2 P$ .

P bits. For small k and P, this is very compact.

## 2. Run-Length Encoding (RLE) for Repeating Patterns:

- **Concept:** Mutations might not be entirely random; they could follow certain patterns. E.g., "for states vx to vy, symbols sa to sb, apply transformation Z." RLE can compress sequences of identical mutation operations.
- **Format:** (START\_V, END\_V, START\_S, END\_S, MUTATION\_TYPE, MUTATION\_PARAM\_1, ...)
- **Size Implications:** Highly dependent on the actual patterns produced by  $\mu$ . Best for structured or "block-wise" mutations. Can lead to variable-size representations, which might require padding to fixed sizes.

## 3. Prefix Maps for Grouped Symbol Transitions:

- **Concept:** Instead of  $\delta(v,s) \rightarrow v'$ , think of groups of symbols. If  $\mu$  affects transitions for a specific range of s from a specific range of v, this can be compressed.
- **Example:** A "prefix map" could define, for v\_prefix || s\_prefix, a specific mutation rule.
- **Format:** A list of (v\_prefix\_mask, s\_prefix\_mask, mutation\_rule\_ID, mutation\_params) tuples.
- **Size Implications:** Reduces overhead if mutations primarily target "blocks" of the state-symbol space.

## 4. Static-Size Representations (64-512 bytes):

- To simplify buffer management and prevent side-channel leakage through size, all mutation packages should be fixed-size.
- This might involve padding smaller compressed results to the max allocated size.
- Example fixed sizes: 64 bytes, 128 bytes, 256 bytes, 512 bytes. The choice depends on the complexity of  $\mu$ 's typical output.

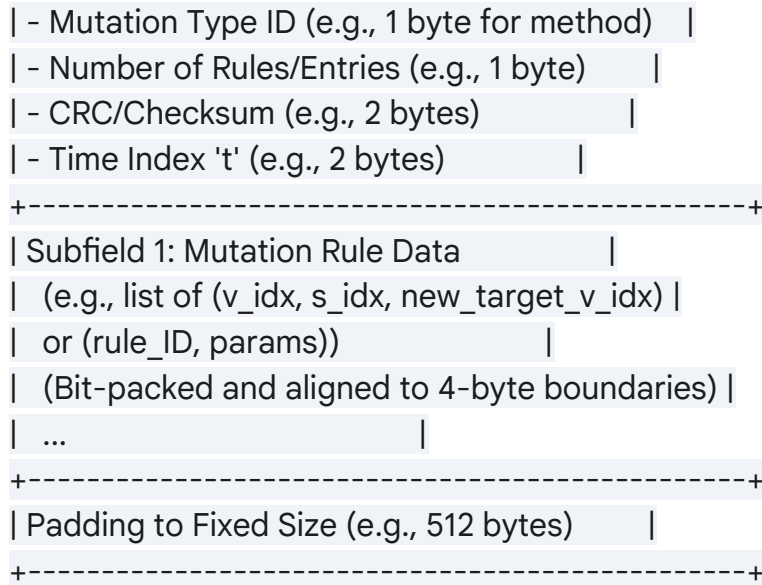
## 3. RAM Buffer Layout for Mutation Execution (RAM Buffer Architect)

The mutation buffer in RAM must be designed for efficiency, constant-time operations, and minimal overhead.

### • Fixed-Size Mutation Buffer ( $\leq 512$ bytes):

- This buffer will hold the compressed mutation result generated by  $\mu(\text{Gbase}, t)$ .
- **Layout:**

```
+-----+
| Header (4-8 bytes) |
```



- **Subfields:**

- **Mutation Index/ID:** A simple integer identifying the specific mutation operation or rule set to be applied. This can be derived from the time index t.
- **Symbol Patch Table / Transition Rewriter:** This is the core data. It could be:
  - A list of (state\_idx, symbol\_idx, new\_target\_state\_idx) tuples, compactly encoded (e.g., 10 bits for state ID, 8 bits for symbol ID, 10 bits for target ID, totaling 28 bits per patch entry).
  - Parameters for a specific rule generator.
  - A small lookup table of permutation indices.
- **Deterministic Hash (e.g., BLAKE2s-32):** A small hash of the mutation package itself, stored in the header. Used by the Reconstruction Verifier to ensure integrity.

- **Minimize Pointer Usage and Stack Depth:**

- Avoid dynamic memory allocation. All data structures are static arrays or pre-allocated fixed buffers.
- Prefer array indexing over pointers where possible to reduce aliasing and potential side-channel risks.
- Strictly limit recursion and function call depth to keep the stack footprint minimal. Iterative algorithms are preferred for applying mutations.

#### 4. Reconstruction Verification (Reconstruction Verifier)

The critical aspect of this approach is ensuring that the compressed mutation data, stored in the low-RAM buffer, can be fully and deterministically reconstructed into the effective Gt for  $\delta$  lookups.

- **Proof of Full and Deterministic Reconstructibility:**

- The mutation\_result (compressed buffer) is a deterministic function of (Gbase, t).
- The reconstruction algorithm Reconstruct(G\_base, mutation\_result) must produce Gt such that for all  $v \in V, s \in \Sigma$ ,  $\delta Gt(v, s) = \delta_{actual\_Gt}(v, s)$ .
- **Formal Argument:** The compression scheme (e.g., delta table, RLE) must be proven lossless. This means that for every unique Gt generated by  $\mu$ , there is a unique and fully recoverable compressed representation, and vice-versa.
- **Mechanism:** The Reconstruct function would typically unpack the compressed patch data and apply it to an in-memory (or Flash-resident) representation of Gbase or a logically empty graph.

- **Proposed Verification Mechanisms:**

1. **Hash Checks (BLAKE2s):**

- During development and testing, after a mutation\_result is generated and compressed, its BLAKE2s hash (e.g., BLAKE2s-32 for 32 bits) is computed and stored.
- After reconstruction: Generate the *full* (uncompressed)  $\delta t$  table. Compute its BLAKE2s hash. Compare it against a pre-computed hash of the expected Gt.
- **Runtime Check:** Store a small hash (e.g., 32-bit CRC or BLAKE2s) of the *compressed* mutation data itself within the mutation buffer header. This allows a quick integrity check on load to ensure the data wasn't corrupted.

2. **Mutation-Seed Replay:**

- The mutation function  $\mu(G_{base}, t)$  itself is deterministic. The most direct way to verify reconstruction is to:
  1. Compute  $G_{t\_true} = \mu(G_{base}, t)$ .
  2. Compress this  $G_{t\_true}$  to mutation\_result\_compressed.
  3. Reconstruct  $G_{t\_reconstructed} = \text{Reconstruct}(G_{base}, \text{mutation\_result\_compressed})$ .
  4. Assert  $G_{t\_true} == G_{t\_reconstructed}$  (i.e., all  $\delta(v, s)$  mappings are identical).
- This is typically done offline during testing, not at runtime on the embedded device.

3. **Incremental Deterministic Generators for Patch Expansion:**

- For some compression schemes (e.g., rule-based), reconstruction might involve generating the full  $\delta t$  table on the fly or sector-by-sector.

- **Verification:** Ensure that the underlying mutation\_rule\_ID corresponds to a unique and deterministic rule\_expansion\_function.
- **Example:** A "rule ID" maps to a specific small function like  $\text{delta\_new}(v, s) = (v \text{ XOR } s) + \text{some\_constant\_from\_params}$ . The verification involves ensuring that all such rule\_expansion\_functions are correct and cover the full range of desired mutations.

## 5. Export Format Management (Meta Role)

Standardizing the mutation package format is crucial for inter-operability, debugging, and formal methods integration.

- **Serialization Format:**

- **Custom Binary Format:** Recommended for embedded runtime mutation packages due to extreme size constraints and parsing efficiency.
  - **Pros:** Minimal overhead, bit-packed, fast parsing.
  - **Cons:** Not human-readable, requires specific parser implementation.
  - **Spec Example:**

```
[0-3]: uint32_t magic_header (e.g., 0xEPMUT)
[4-5]: uint16_t version (e.g., 0x0100)
[6-7]: uint16_t mutation_type_id
[8-9]: uint16_t time_index_t
[10-11]: uint16_t num_patch_entries
[12-15]: uint32_t crc32_checksum // or BLAKE2s-32 hash
[16-N]: byte[] compressed_patch_data // Bit-packed as per
Symbol_Compression_Engineer's design
[N+1-M]: byte[] padding_to_fixed_size // e.g., 512 bytes
```
- **JSON/CBOR:** Preferred for debug/test mutation logs and formal proof simulation inputs.
  - **JSON:** Human-readable, easy for scripting, but verbose. Good for initial prototyping.
  - **CBOR (Concise Binary Object Representation):** Binary-encoded JSON, more compact than JSON but still self-describing. Good for test vectors exchange.
  - **Spec Example (JSON):**

```
JSON
{
  "magic_header": "EPMUT",
```

```

"version": "1.0",
"mutation_type": "delta_remap",
"time_index": 1234,
"checksum_blake2s": "abcdef1234567890", // Hash of 'patch_data'
"patch_data": [
  {"v_start": 0, "s_val": 1, "v_target": 12},
  {"v_start": 5, "s_val": 255, "v_target": 1023},
  // ... compressed representation of changes
]
}

```

- **Validation Rules for Formats:**

- **Schema Definition:** For JSON/CBOR, provide a JSON Schema to ensure correct structure.
- **Checksum Verification:** Mandate the use of a checksum (e.g., CRC32 or BLAKE2s) in the header of all mutation packages to detect transmission or storage errors.
- **Range Checks:** Ensure all indices and values within the compressed patch data fall within valid ranges (e.g.,  $v \in [0, |V| - 1]$ ,  $s \in [0, |\Sigma| - 1]$ ).

## 6. Final Recommendations for Integrators

- **Prioritize Bit-Packing:** For the actual compressed\_patch\_data, bit-packing and custom encoding will yield the most significant RAM savings.
- **Fixed Size is Key:** Design all mutation buffers and packages to be fixed sizes to eliminate timing side-channels related to data length. Pad smaller packages if necessary.
- **Offline Compression & Verification:** The most complex compression and reconstruction verification (e.g.,  $G_t\_true == G_t\_reconstructed$ ) should happen offline during development and testing. The embedded device only performs the simpler runtime reconstruction.
- **Test with Worst-Case Mutations:** Benchmark performance and memory usage with mutation patterns that are least compressible to ensure the chosen fixed buffer size is sufficient for all t.
- **Integrate with Low-RAM Operating Mode:** This mutation compression scheme is a vital component of the overall "Low-RAM Operating Mode" for EchoPulse. Both designs must be implemented in conjunction.

This design document provides a robust framework for managing and optimizing the

RAM footprint of EchoPulse's dynamic mutation, enabling its deployment on the most memory-constrained platforms while upholding its cryptographic integrity.